

Synchronization

Learning Outcomes

- ♦ Data Race, Mutual Exclusion, Deadlocks
- ♦ Atomics, Locks, Barriers
- ♦ Reduction
- ♦ Prefix Sum

Classwork

- Write CUDA code for the following functionality.
 - Assume following data type, filled with some values.
`struct Point { int x, y; } arr[N];`
 - Each thread should operate on 4 elements.
 - Find the average AVG of x values.
 - If a thread sees y value above the average, it replaces all 4 y values with AVG.
 - Otherwise, it adds y values to a global sum.
 - Host prints the number of elements set to AVG.

Data Race

- A data race occurs if *all* of the following hold:
 1. Multiple threads
 2. Common memory location
 3. At least one write
 4. Concurrent execution
- Ways to remove data race:
 1. Execute sequentially
 2. Privatization / Data replication
 3. Separating reads and writes by a barrier
 4. Mutual exclusion

Classwork


- Is there a data race in this code?
- What does the code ensure?
- Can mutual exclusion be generalized for N threads?

If initially `flag == 0`, then S2 executes before S1.
If initially `flag == 1`, then S2 executes and after that S1 may execute or T1 may hang.

T1	T2
<pre>flag = 1; while (flag) ; S1;</pre>	<pre>while (!flag) ; S2; flag = 0;</pre>


Classwork: Grading

- Given roll numbers and marks of 80 students in GPU Programming, assign grades.
 - S = 90, A = 80, B = 70, ..., E = 40, and U.
 - Use input arrays and output arrays.
- Compute the histogram.
 - Count the number of students with a grade.




Let's Compute the Shortest Paths

- You are given an input graph of India, and you want to compute the shortest path from Nagpur to every other city.
- Assume that you are given a GPU graph library and the associated routines.
- Each thread operates on a node and settles distances of the neighbors (Bellman-Ford style).



Let's Compute the Shortest Paths

- You are given an input graph of India, and you want to compute the shortest path from Nagpur to every other city.
- Assume that you are given a GPU graph library and the associated routines.



```
__global__ void dsssp(Graph g, unsigned *dist) {  
    unsigned id = ...  
    for each n in g.allneighbors(id) {    // pseudo-code.  
        unsigned altdist = dist[id] + weight(id, n);  
        if (altdist < dist[n]) {  
            dist[n] = altdist;  
        }  
    }  
}
```

What is the error in this code?