

# Synchronization


# Recap - atomics

```
__global__ void dkernel(int *x) {  
    ++x[0];  
}  
...  
dkernel<<<2, 1>>>(x);
```

- Ensure all-or-none behavior.
  - e.g., `atomicInc(&x[0], ...);`
- `dkernel<<<K1, K2>>>` would ensure `x[0]` to be incremented by exactly  $K1 * K2$  – irrespective of the thread execution order.
  - When would this effect be visible?

# Let's Compute the Shortest Paths


- You are given an input graph of India, and you want to compute the shortest path from Nagpur to every other city.
- Assume that you are given a GPU graph library and the associated routines.



```
__global__ void dsssp(Graph g, unsigned *dist) {  
    unsigned id = ...  
    for each n in g.allneighbors(id) {    // pseudo-code.  
        unsigned altdist = dist[id] + weight(id, n);  
        if (altdist < dist[n]) {  
            dist[n] = altdist;  
        }  
    }  
}
```

# Let's Compute the Shortest Paths

- You are given an input graph of India, and you want to compute the shortest path from Nagpur to every other city.
- Assume that you are given a GPU graph library and the associated routines.



```
__global__ void dsssp(Graph g, unsigned *dist) {
    unsigned id = ...
    for each n in g.allneighbors(id) {    // pseudo-code.
        unsigned altdist = dist[id] + weight(id, n);
        if (altdist < dist[n]) {
            dist[n] = altdist;    atomicMin(&dist[n], altdist);
        }
    }
}
```

# AtomicCAS

- Syntax: `oldval = atomicCAS(&var, x, y);`
- **Typical usecases:**
  - *Locks*: critical section processing
  - *Single*: Only one arbitrary thread executes the block.

**Classwork: Implement *lock* with *atomicCAS*.**

# Lock using atomicCAS

Does this work?

```
atomicCAS(&lockvar, 0, 1);
```

Does not ensure mutual exclusion.

Then how about

```
if (atomicCAS(&lockvar, 0, 1) == 0)
```

```
    // critical section
```

Does not block other threads.

Make the above code blocking.

```
do {
```

```
    old = atomicCAS(&lockvar, 0, 1);
```

```
} while (old != 0);
```

Correct code?

# Lock using atomicCAS

- The code works on CPU.
- It also works on GPU across warps.
- But it hangs for threads belonging to the same warp.
  - When one warp-thread acquires the lock, it waits for other warp-threads to reach the instruction just after the **do-while**.
  - Other warp-threads await this successful thread in the **do-while**.

```
do {  
  
    old = atomicCAS(&lockvar, 0, 1);  
  
} while (old != 0);
```



Correct code?

# Lock using atomicCAS

```
do {  
    old = atomicCAS(&lockvar, 0, 1);  
} while (old != 0);
```

**// critical section**

lockvar =

On CPU

```
do {  
    old = atomicCAS(&lockvar, 0, 1);  
    if (old == 0) {  
        // critical section  
        lockvar = 0; // unlock  
    }  
} while (old != 0);
```

On GPU

**Classwork: Implement *single* with *atomicCAS*.**



# Single using atomicCAS

```
if (atomicCAS(&lockvar, 0, 1) == 0)
```

```
    // single section
```

Important not to set lockvar to 0 at the end of the single section.

# What is the output?

```
#include <stdio.h>
```

```
#include <cuda.h>
```

```
__global__ void k1(int *gg) {  
    int old = atomicCAS(gg, 0, threadIdx.x + 1);  
    if (old == 0) {  
        printf("Thread %d succeeded 1.\n", threadIdx.x);  
    }  
    old = atomicCAS(gg, 0, threadIdx.x + 1);  
    if (old == 0) {  
        printf("Thread %d succeeded 2.\n", threadIdx.x);  
    }  
    old = atomicCAS(gg, threadIdx.x, -1);  
    if (old == threadIdx.x) {  
        printf("Thread %d succeeded 3.\n", threadIdx.x);  
    }  
}  
  
int main() {  
    int *gg;  
    cudaMalloc(&gg, sizeof(int));  
    cudaMemset(&gg, 0, sizeof(int));  
    k1<<<2, 32>>>(gg);  
    cudaDeviceSynchronize();  
  
    return 0;  
}
```

- Some thread out of 64 updates gg to its threadid+1.
- Warp threads do not execute atomics together! That is also done sequentially.
- Irrespective of which thread executes the first atomicCAS, no thread would see gg to be 0. Hence second printf is not executed at all.
- If gg was updated by some thread 0..30, then the corresponding thread with id 1..31 from either of the blocks would update gg to -1, and execute the third printf.
- Otherwise, no one would update gg to -1, and no one would execute the third printf.
- On most executions, you would see the output to be that thread 0 would execute the first printf, and thread 1 would execute the third printf.

# Synchronization

- Control + data flow
- Atomics
- Barriers
- ...

# Barriers

- A barrier is a program point where all threads need to reach before any thread can proceed.
- End of kernel is an implicit barrier for all GPU threads (**global barrier**).
- ~~There is no explicit global barrier supported in CUDA.~~ *grid.sync()* is now supported (from CUDA 9).
- Threads in a thread-block can synchronize using **\_\_syncthreads()**.
- How about barrier within warp-threads?