

Functions

CUDA Function Declarations

	Executed on the:	Callable from only the:
<code>__device__ float DeviceFunc()</code>	device	device
<code>__global__ void KernelFunc()</code>	device	host
<code>__host__ float HostFunc()</code>	host	host

- `__global__` defines a kernel. It must return `void`.
- A program may have several functions of each kind.
- The same function of any kind may be called multiple times.
- Host == CPU, Device == GPU.

Function Types (1/2)


```
#include <stdio.h>
#include <cuda.h>
__host__ __device__ void dhfun() {
    printf("I can run on both CPU and GPU.\n");
}
__device__ unsigned dfun(unsigned *vector, unsigned vectorsize, unsigned id) {
    if (id == 0) dhfun();
    if (id < vectorsize) {
        vector[id] = id;
        return 1;
    } else {
        return 0;
    }
}
__global__ void dkernel(unsigned *vector, unsigned vectorsize) {
    unsigned id = blockIdx.x * blockDim.x + threadIdx.x;
    dfun(vector, vectorsize, id);
}
__host__ void hostfun() {
    printf("I am simply like another function running on CPU. Calling dhfun\n");
    dhfun();
}
```

Function Types (2/2)

```
#define BLOCKSIZE    1024
int main(int nn, char *str[]) {
    unsigned N = atoi(str[1]);
    unsigned *vector, *hvector;
    cudaMalloc(&vector, N * sizeof(unsigned));
    hvector = (unsigned *)malloc(N * sizeof(unsigned));

    unsigned nblocks = ceil((float)N / BLOCKSIZE);
    printf("nblocks = %d\n", nblocks);

    dkernel<<<nblocks, BLOCKSIZE>>>(vector, N);
    cudaMemcpy(hvector, vector, N * sizeof(unsigned), cudaMemcpyDeviceToHost);
    for (unsigned ii = 0; ii < N; ++ii) {
        printf("%4d ", hvector[ii]);
    }
    printf("\n");
    hostfun();
    dhfun();
    return 0;
}
```



What are the other arrows possible in this diagram?
How about **dhfun** to **dfun**?

with HostAlloc'ed Memory

`__host__ __device__`
functions are friends with
`HostAlloc`'ed memory.

```
__host__ __device__ void fun(int *counter) {  
    ++*counter;  
}  
__global__ void printk(int *counter) {  
    fun(counter);  
    printf("printk (after fun): %d\n", *counter);  
}  
int main() {  
    int *counter;  
    cudaHostAlloc(&counter, sizeof(int), 0);  
  
    *counter = 0;  
    printf("main: %d\n", *counter);  
  
    printk<<<1, 1>>>(counter);  
    cudaDeviceSynchronize();  
  
    fun(counter);  
    printf("main (after fun): %d\n", *counter);  
  
    return 0;  
}
```

What is the
output of
this code?

with a Device-only Function

```
__host__ __device__ void fun(int *counter) {  
    ++*counter;  
    __syncthreads();  
}  
__global__ void printk(int *counter) {  
    fun(counter);  
    printf("printk (after fun): %d\n", *counter);  
}  
int main() {  
    int *counter;  
    cudaHostAlloc(&counter, sizeof(int), 0);  
  
    *counter = 0;  
    printf("main: %d\n", *counter);  
  
    printk <<<1, 1>>>(counter);  
    cudaDeviceSynchronize();  
  
    fun(counter);  
    printf("main (after fun): %d\n", *counter);  
  
    return 0;  
}
```

__syncthreads()
is not available
on CPU.

with a CPU-only Memory

```
__host__ __device__ void fun(int *counter) {  
    ++*counter;  
}  
  
__global__ void printk(int *counter) {  
    fun(counter);  
    printf("printk (after fun): %d\n", *counter);  
}  
  
int main() {  
    int *counter;  
    // cudaHostAlloc(&counter, sizeof(int), 0);  
    cudaMalloc(&counter, sizeof(int));  
  
    *counter = 0;  
    printf("main: %d\n", *counter);  
  
    printk <<<1, 1>>>(counter);  
    cudaDeviceSynchronize();  
  
    fun(counter);  
    printf("main (after fun): %d\n", *counter);  
  
    return 0;  
}
```

counter cannot
be accessed
on CPU.

Global Variables

```
int counter;

__host__ __device__ void fun() {
    ++counter;
}

__global__ void printk() {
    fun();
    printf("printk (after fun): %d\n", counter);
}

int main() {

    counter = 0;
    printf("main: %d\n", counter);

    printk <<<1, 1>>>();
    cudaDeviceSynchronize();

    fun();
    printf("main (after fun): %d\n", counter);

    return 0;
}
```

counter cannot
be accessed
on **GPU**.

Global Variables

```
__host__ __device__ int counter;

__host__ __device__ void fun() {
    ++counter;
}

__global__ void printk() {
    fun();
    printf("printk (after fun): %d\n", counter);
}

int main() {

    counter = 0;
    printf("main: %d\n", counter);

    printk <<<1, 1>>>();
    cudaDeviceSynchronize();

    fun();
    printf("main (after fun): %d\n", counter);

    return 0;
}
```

Variables
cannot be
declared as
__host__.

Global Variables

```
__device__ int counter;

__host__ __device__ void fun() {
    ++counter;
}

__global__ void printk() {
    fun();
    printf("printk (after fun): %d\n", counter);
}

int main() {

    printk <<<1, 1>>>();
    cudaDeviceSynchronize();

    return 0;
}
```

Warning during compilation, but works fine.

Classwork

Write a CUDA code to increment all elements in an array. Call this code from host as well as device.

Classwork: Can you avoid the for loop in **fun**?

```
__host__ __device__ void fun(int *arr) {  
    for (unsigned ii = 0; ii < N; ++ii)  
        ++arr[ii];  
}  
__global__ void dfun(int *arr) {  
    fun(arr);  
}  
int main() {  
    int arr[N], *darr;  
  
    cudaMalloc(&darr, N * sizeof(int));  
  
    for (unsigned ii = 0; ii < N; ++ii)  
        arr[ii] = ii;  
    cudaMemcpy(darr, arr, N * sizeof(int),  
               cudaMemcpyHostToDevice);  
  
    fun(arr);  
    dfun<<<1, 1>>>(darr);  
    cudaDeviceSynchronize();  
  
    return 0;  
}
```

Host-centric,
sequential on GPU

Classwork

Write a CUDA code to increment all elements in an array. Call this code from host as well as device.

Classwork: Can you avoid the for loop in **fun**?

Classwork: What if I don't like the for loop in main, but still want GPU-parallel code?

```
__host__ __device__ void fun(int *arr) {  
    ++(*arr);  
}  
__global__ void dfun(int *arr) {  
    fun(arr + threadIdx.x);  
}  
int main() {  
    int arr[N], *darr;  
  
    cudaMalloc(&darr, N * sizeof(int));  
  
    for (unsigned ii = 0; ii < N; ++ii)  
        arr[ii] = ii;  
    cudaMemcpy(darr, arr, N * sizeof(int),  
               cudaMemcpyHostToDevice);  
    for (unsigned ii = 0; ii < N; ++ii)  
        fun(arr + ii);  
    dfun<<<1, N>>>>(darr);  
    cudaDeviceSynchronize();  
  
    return 0;  
}
```

Device-centric,
sequential on CPU

Classwork

Write a CUDA code to increment all elements in an array. Call this code from host as well as device.

Classwork: Can you avoid the for loop in **fun**?

Classwork: What if I don't like the for loop in main, but still want GPU-parallel code?

```
__host__ __device__ void fun(int *arr, int nn) {
    for (unsigned ii = 0; ii < nn; ++ii)
        ++arr[ii];
}

__global__ void dfun(int *arr) {
    fun(arr + threadIdx.x, 1);
    // need to change for more blocks.
}

int main() {
    int arr[N], *darr;

    cudaMalloc(&darr, N * sizeof(int));

    for (unsigned ii = 0; ii < N; ++ii)
        arr[ii] = ii;
    cudaMemcpy(darr, arr, N * sizeof(int),
               cudaMemcpyHostToDevice);

    fun(arr, N);
    dfun<<<1, N>>>(darr);
    cudaDeviceSynchronize();

    return 0;
}
```

Thrust

- Thrust is a parallel algorithms library (similar in spirit to STL on CPU).
- Supports vectors and associated transforms.
- Programmer is oblivious to where code executes
 - on CPU or GPU.