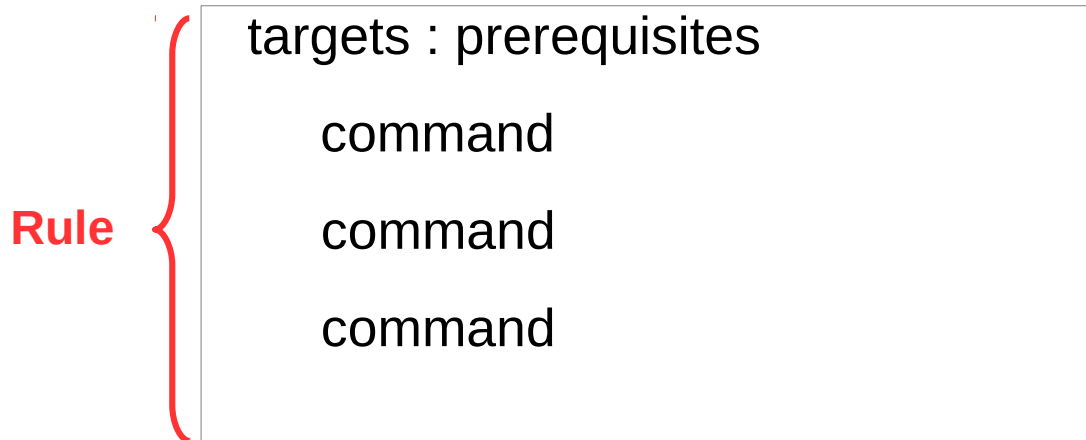


Support: Makefiles, Debugging and Profiling

Makefile

- Makefile is a set of commands
- Makefiles are used when we have dependencies among files
- It is helpful for the large project where there might be many dependencies among the files
- make utility is a command line utility used to process the instructions in the Makefile
- Put the commands in a file called Makefile, and in that directory run the command make

Makefile Syntax



- The targets and prerequisites are file names or actions, separated by colon
- The commands are a series of steps
- Note: commands need to start with a tab character, not spaces

Example

Hello:

echo "hello world"

```
$ make  
echo "hello world"  
hello world
```

Running CUDA Program

```
#include <stdio.h>
#include <cuda.h>
__global__ void dkernel() {
    printf("Hello World.\n");
}
int main() {
    dkernel<<<1, 32>>>();
    cudaDeviceSynchronize();
    return 0;
}
```

32 times {

```
clean : blah
    rm hello

blah : hello
    ./hello

hello :
    nvcc hello.cu -o hello
```

```
$ make
Hello World.
Hello World.
...
```

Variables

- Variables can only be strings
- Typically you can use `:=` for assignment
- Reference variables using either `${}` or `$()`

```
file1 := hello
```

```
file2 := blah
```

```
file3 := clean
```

```
${file3} : $(file2)
```

```
rm hello
```

```
${file2} : $(file1)
```

```
./hello
```

```
$(file1) :
```

```
nvcc hello.cu -o hello
```

```
$ make
```

```
Hello World.
```

```
Hello World.
```

```
...
```

All target

→ We can use all target for making multiple targets and run them all.

```
all: H1 H2 H3
```

```
H1:
```

```
    nvcc hello1.cu -o hello1
```

```
    ./hello1
```

```
H2:
```

```
    nvcc hello2.cu -o hello2
```

```
    ./hello2
```

```
H3:
```

```
    nvcc hello3.cu -o hello3
```

```
    ./hello3
```

```
clean:
```

```
    rm -f hello1 hello2 hello3
```

Multiple target

→ When there are multiple targets for a rule, the commands will be run for each target

```
all: call1 call2 call3
```

```
call1 call2 call3:
```

```
    nvcc algo1.cu -o algo1
```

```
    ./algo1
```

```
    nvcc algo2.cu -o algo2
```

```
    ./algo2
```

```
clean:
```

```
    rm -f algo1 algo2
```


Conditions in Makefiles

```
foo = somestring
all:
ifneq ($(foo), ok)
    echo "foo not equals ok"
else
    echo "nope"
endif
```

```
TARGET_CPU_IS_X86 := 1
all:
ifeq ( $(TARGET_CPU), x86 )
    TARGET_CPU_IS_X86 := 1
else ifeq( $(TARGET_CPU), x86_64 )
    TARGET_CPU_IS_X86 := 1
else
    TARGET_CPU_IS_X86 := 0
endif
```

Loops in Makefiles

all:

nvcc algo.cu -o algo

for number in 32 100 900 ; do \

./algo \$\$number ; \

done

all:

nvcc algo.cu -o algo

number=100 ; while [[\$\$number -le
1000]] ; do \

./algo \$\$number ; \

((number = number + 100)) ; \

done

Example

```
all : vertex-removal

    unzip examples/test_cases/Bench.zip;      mv examples/test_cases/*.txt ./ ;
    number=1 ;
    while [[ $$number -le 10 ]] ; do \
        ./vertex-removal t$$number.mtx > vr_graph_$$number.txt ; \
        ifeq ( $$number , 4 )
            ((number = number + 2)) ;\
        else
            ((number = number + 1)) ;\
        endif
    done

clean : all

    rm t*.mtx; rm vertex-removal;

vertex-removal:

    nvcc -std=c++11 -o vertex-removal vertex-removal.cu;
```

Debugging

- Debugging parallel programs is difficult.
 - Non-determinism due to thread-scheduling
 - Output can be different
 - Correct intermediate values may be large
- cuda-gdb
 - for debugging CUDA programs on real hardware
 - Extension to gdb
 - Allows breakpoints, single-step, read/write memory contents.

Sample Error

```
#include <cuda.h>
#include <stdio.h>
```

```
__global__ void K(int *x) {
    *x = 0;
}
int main() {
    int *x;
    K<<<2, 10>>>(x);
    cudaDeviceSynchronize();

    return 0;
}
```

Sample Error

```
#include <cuda.h>
```

```
#include <stdio.h>
```

```
__global__ void K(int *x) {
```

```
    *x = 0;
```

```
    printf("%d\n", *x);    // does not print anything.
```

```
}
```

```
int main() {
```

```
    int *x;
```

```
    K<<<2, 10>>>(x);
```

```
    cudaDeviceSynchronize();
```

```
    return 0;
```

```
}
```

Sample Error

```
#include <cuda.h>
#include <stdio.h>

__global__ void K(int *x) {
    *x = 0;
    printf("%d\n", *x);
}

int main() {
    int *x;
    K<<<2, 10>>>(x);
    cudaDeviceSynchronize();
    cudaError_t err = cudaGetLastError();
    printf("error=%d, %s, %s\n", err, cudaGetErrorName(err),
        cudaGetErrorString(err));
    return 0;
}
```

**error=77, cudaErrorIllegalAddress,
an illegal memory access was encountered**

CUDA Errors

- `cudaSuccess` = 0, */// No errors*
- `cudaErrorMissingConfiguration` = 1, */// Missing configuration error*
- `cudaErrorMemoryAllocation` = 2, */// Memory allocation error*
- `cudaErrorInitializationError` = 3, */// Initialization error*
- `cudaErrorLaunchFailure` = 4, */// Launch failure*
- `cudaErrorPriorLaunchFailure` = 5, */// Prior launch failure*
- `cudaErrorLaunchTimeout` = 6, */// Launch timeout error*
- `cudaErrorLaunchOutOfResources` = 7, */// Launch out of resources*
- `cudaErrorInvalidDeviceFunction` = 8, */// Invalid device function*
- `cudaErrorInvalidConfiguration` = 9, */// Invalid configuration*
- `cudaErrorInvalidDevice` = 10, */// Invalid device*
- ...

Homework: Write programs to invoke these errors.

cuda-gdb

- Generate debug information
 - `nvcc -g -G file.cu`
 - Disables optimizations.
- Run with cuda-gdb
 - `cuda-gdb a.out`
 - `> run`
- Due to lots of threads, cuda-gdb works with a focus (current thread).

```

__global__ void K(int *x)
{
    *x = 0;
    printf("%d\n", *x);
}

```

(cuda-gdb) **run**

Starting program:a.out

[Thread debugging using libthread_db enabled]

Using host libthread_db library "/lib64/libthread_db.so.1".

[New Thread 0x7fff7396700 (LWP 10305)]

[New Thread 0x7fff696d700 (LWP 10306)]

CUDA Exception: Device Illegal Address

The exception was triggered in device 0.

Program received signal CUDA_EXCEPTION_10, Device Illegal Address.

[Switching focus to CUDA kernel 0, grid 1, block (1,0,0), thread (0,0,0), device 0, sm 13, warp 0, lane 0]

0x0000000000aa9510 in K<<<(2,1,1),(10,1,1)>>> (x=0x0) at gdb2.cu:6

6 printf("%d\n", *x);

(cuda-gdb)

cuda-gdb

(cuda-gdb) **info cuda kernels**

	Kernel	Parent	Dev	Grid	Status	SMs	Mask	GridDim	BlockDim	Invocation
*	0	-	0	1	Active	0x00006000		(2,1,1)	(10,1,1)	K(x=0x0)

(cuda-gdb) **info threads**

	Id	Target	Id	Frame
	3	Thread	0x7fff696d700	(LWP 10497) "a.out" 0x00000038db4df113 in poll () from /lib64/libc.so.6
	2	Thread	0x7fff7396700	(LWP 10496) "a.out" 0x00000038db4eac6f in accept4 () from /lib64/libc.so.6
*	1	Thread	0x7fff7fca720	(LWP 10487) "a.out" 0x00007ffff77a2118 in cudbgApiDetach () from /usr/lib64/libcuda.so.1

cuda-gdb

(cuda-gdb) **info cuda threads**

BlockIdx	ThreadIdx	To BlockIdx	ThreadIdx	Count	Virtual PC	Filename	Line
Kernel 0							
*	(0,0,0)	(0,0,0)	(1,0,0)	(9,0,0)	20	0x000000000000aa9f50	gdb2.cu 6

(cuda-gdb) **cuda kernel block thread**

kernel 0, block (0,0,0), thread (0,0,0)

(cuda-gdb) **cuda block 1 thread 0**

[Switching focus to CUDA kernel 0, grid 1, block (1,0,0), thread (0,0,0), device 0, sm 13, warp 0, lane 0]

0x000000000000aa9510 6 printf("%d\n", *x);

(cuda-gdb) **cuda kernel block thread**

kernel 0, block (**1**,0,0), thread (0,0,0)

Breakpoints

- `break main` // first instruction in main
- `break file.cu:223` // file:line
- `set cuda break_on_launch application`
// kernel entry breakpoint
- `break file.cu:23 if threadIdx.x == 1 && i < 5`
// conditional breakpoint

Step

- Once at a breakpoint, you can single-step
 - step, s or <enter>

(cuda-gdb) **info cuda sms**

SM Active Warps Mask

Device 0

0 0x00000000000000000000

1 0x00000000000000000000

2 0x00000000000000000000

3 0x00000000000000000000

4 0x00000000000000000000

5 0x00000000000000000000

6 0x00000000000000000000

7 0x00000000000000000000

8 0x00000000000000000000

9 0x00000000000000000000

10 0x00000000000000000000

11 0x00000000000000000000

12 0x00000000000000000000

13 0x00000000000000000000**1**

* 14 0x00000000000000000000**1**

(cuda-gdb) **info cuda warps**

Wp Active Lanes Mask Divergent Lanes Mask Active Physical PC Kernel BlockIdx First Active ThreadIdx

Device 0 SM 14

* 0	0x000003ff	0x00000000	0x0000000000000000110	0	(0,0,0)	(0,0,0)
1	0x00000000	0x00000000	n/a	n/a	n/a	n/a
2	0x00000000	0x00000000	n/a	n/a	n/a	n/a
3	0x00000000	0x00000000	n/a	n/a	n/a	n/a
4	0x00000000	0x00000000	n/a	n/a	n/a	n/a
5	0x00000000	0x00000000	n/a	n/a	n/a	n/a
6	0x00000000	0x00000000	n/a	n/a	n/a	n/a
7	0x00000000	0x00000000	n/a	n/a	n/a	n/a
8	0x00000000	0x00000000	n/a	n/a	n/a	n/a
9	0x00000000	0x00000000	n/a	n/a	n/a	n/a
10	0x00000000	0x00000000	n/a	n/a	n/a	n/a

---Type <return> to continue, or q <return> to quit---

(cuda-gdb) **info cuda lanes**

Ln	State	Physical PC	ThreadIdx	Exception
Device 0 SM 14 Warp 0				
* 0	active	0x0000000000000000110	(0,0,0)	Device Illegal Address
1	active	0x0000000000000000110	(1,0,0)	Device Illegal Address
2	active	0x0000000000000000110	(2,0,0)	Device Illegal Address
3	active	0x0000000000000000110	(3,0,0)	Device Illegal Address
4	active	0x0000000000000000110	(4,0,0)	Device Illegal Address
5	active	0x0000000000000000110	(5,0,0)	Device Illegal Address
6	active	0x0000000000000000110	(6,0,0)	Device Illegal Address
7	active	0x0000000000000000110	(7,0,0)	Device Illegal Address
8	active	0x0000000000000000110	(8,0,0)	Device Illegal Address
9	active	0x0000000000000000110	(9,0,0)	Device Illegal Address
10	inactive	n/a	n/a	n/a
11	inactive	n/a	n/a	n/a
12	inactive	n/a	n/a	n/a
...				
29	inactive	n/a	n/a	n/a
30	inactive	n/a	n/a	n/a
31	inactive	n/a	n/a	n/a

Homework

For the given program, what sequence of cuda-gdb commands would you use to identify the error?

```
__global__ void K(int *p) {  
    *p = 0;  
    printf("%d\n", *p);  
}  
int main() {  
    int *x, *y;  
    cudaMalloc(&x, sizeof(int));  
    K<<<2, 10>>>(x);  
    cudaDeviceSynchronize();  
    y = x;  
    cudaFree(y);  
    K<<<2, 10>>>(x);  
    cudaDeviceSynchronize();  
    return 0;  
}
```

Profiling

- Measuring “indicators” of performance
 - Time taken by various kernels
 - Memory utilization
 - Number of cache misses
 - Degree of divergence
 - Degree of coalescing
 - ...
- Intrusive versus non-intrusive

CUDA Profiler

- **nvprof**: command-line
- **nvvp**, **nsight**: Visual Profilers

nvprof

- No changes required to the binary. Uses defaults.
 - nvprof a.out
- To profile part of a program, use *cudaProfilerStart()* and *Stop()*.
 - Include *cuda_profiler_api.h*
 - nvprof **--profile-from-start off** a.out

```

__global__ void K1(int num) {
    num += num;
    ++num;
}

__device__ int sum = 0;
__global__ void K2(int num) {
    atomicAdd(&sum, num);
}

__global__ void K3(int num) {
    __shared__ int sum;
    sum = 0;
    __syncthreads();
    sum += num;
}

int main() {
    for (unsigned ii = 0; ii < 100; ++ii) {
        K1<<<5, 32>>>(ii);        cudaDeviceSynchronize();
    }
    for (unsigned ii = 0; ii < 100; ++ii) {
        K2<<<5, 32>>>(ii);        cudaDeviceSynchronize();
    }
    for (unsigned ii = 0; ii < 100; ++ii) {
        K3<<<5, 32>>>(ii);        cudaDeviceSynchronize();
    }
    return 0;
}

```

Which kernel should you optimize?
 (Which kernel consumes more time?)

\$ nvprof a.out

==26519== NVPROF is profiling process 26519, command: a.out

==26519== Profiling application: a.out

==26519== Profiling result:

Time(%)	Time	Calls	Avg	Min	Max	Name
39.46%	191.46us	100	1.9140us	1.8880us	2.1440us	K2(int)
33.86%	164.26us	100	1.6420us	1.6000us	1.8880us	K3(int)
26.68%	129.44us	100	1.2940us	1.2480us	1.5360us	K1(int)

==26519== API calls:

Time(%)	Time	Calls	Avg	Min	Max	Name
95.75%	369.08ms	300	1.2303ms	10.560us	364.03ms	cudaLaunch
2.33%	8.9986ms	728	12.360us	186ns	619.78us	cuDeviceGetAttribute
0.91%	3.5039ms	8	437.98us	396.85us	450.61us	cuDeviceTotalMem
0.73%	2.8134ms	300	9.3780us	6.4650us	32.547us	cudaDeviceSynchronize
0.18%	699.99us	8	87.498us	85.431us	90.737us	cuDeviceGetName
0.05%	194.20us	300	647ns	339ns	10.694us	cudaConfigureCall
0.04%	156.27us	300	520ns	292ns	2.2700us	cudaSetupArgument
0.00%	9.4130us	24	392ns	186ns	862ns	cuDeviceGet
0.00%	5.7760us	3	1.9250us	317ns	4.7490us	cuDeviceGetCount

nvprof

- Supports device-specific profiling
- Supports remote profiling
- Output can be dumped to files as a .csv
- ...