

# Topics

# Topics

- Dynamic Parallelism
- Multi-GPU Processing
- Warp Voting

# Dynamic Parallelism

- Useful in scenarios involving nested parallelism.

```
for i ...  
  for j = f(i) ...  
    work(j)
```

- Algorithms using hierarchical data structures
  - Algorithms using recursion where each level of recursion has parallelism
  - Algorithms where work naturally splits into independent batches, and each batch involves parallel processing
- Not all nested parallel loops need DP.

```

#include <stdio.h>
#include <cuda.h>
__global__ void Child(int father) {
    printf("Parent %d -- Child %d\n", father, threadIdx.x);
}
__global__ void Parent() {
    printf("Parent %d\n", threadIdx.x);
    Child<<<1, 5>>>(threadIdx.x);
}
int main() {
    Parent<<<1, 3>>>();
    cudaDeviceSynchronize();
    return 0;
}

```

**\$ nvcc dynpar.cu**

**error:** calling a \_\_global\_\_ function("Child") from a \_\_global\_\_ function("Parent") is only allowed on the compute\_35 architecture or above

**\$ nvcc -arch=sm\_35 dynpar.cu**

**error:** kernel launch from \_\_device\_\_ or \_\_global\_\_ functions requires separate compilation mode

**\$ nvcc -arch=sm\_35 -rdc=true dynpar.cu**

**\$ a.out**

```

#include <stdio.h>
#include <cuda.h>
__global__ void Child(int father) {
    printf("Parent %d -- Child %d\n", father, threadIdx.x);
}
__global__ void Parent() {
    printf("Parent %d\n", threadIdx.x);
    Child<<<1, 5>>>(threadIdx.x);
}
int main() {
    Parent<<<1, 3>>>();
    cudaDeviceSynchronize();
    return 0;
}

```

**Parent 0**  
**Parent 1**  
**Parent 2**  
**Parent 0 -- Child 0**  
**Parent 0 -- Child 1**  
**Parent 0 -- Child 2**  
**Parent 0 -- Child 3**  
**Parent 0 -- Child 4**  
**Parent 1 -- Child 0**  
**Parent 1 -- Child 1**  
**Parent 1 -- Child 2**  
**Parent 1 -- Child 3**  
**Parent 1 -- Child 4**  
**Parent 2 -- Child 0**  
**Parent 2 -- Child 1**  
**Parent 2 -- Child 2**  
**Parent 2 -- Child 3**  
**Parent 2 -- Child 4**

```
#include <stdio.h>
#include <cuda.h>
```

```
#define K 2
```

```
__global__ void Child(int father) {
    printf("%d\n", father + threadIdx.x);
}
```

```
__global__ void Parent() {
    if (threadIdx.x % K == 0) {
        Child<<<1, K>>>(threadIdx.x);
        printf("Called children with starting %d\n", threadIdx.x);
    }
}
```

```
int main() {
    Parent<<<1, 10>>>();
    cudaDeviceSynchronize();

    return 0;
}
```

```
0
1
Called children with starting 0
Called children with starting 2
Called children with starting 4
Called children with starting 6
Called children with starting 8
2
3
4
5
6
7
8
9
```

# DP: Computation

- Parent kernel is associated with a parent grid.
- Child kernels are associated with child grids.
- Parent and child kernels may execute asynchronously.
- A parent grid is not complete unless all its children have completed.

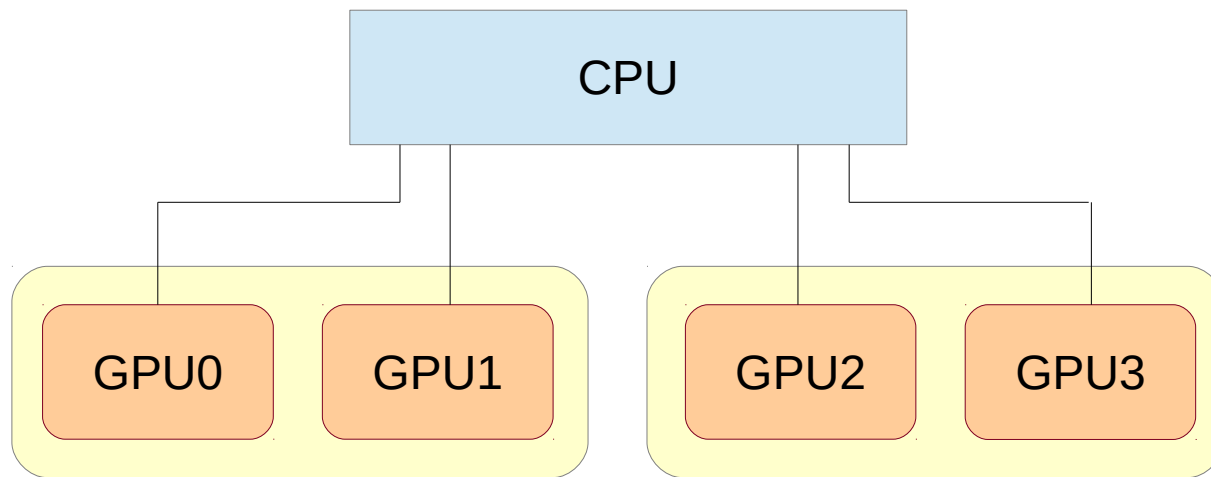
# DP: Memory

- Parent and children **share** global and constant memory.
- But they have **distinct** local and shared memories.
- All global memory operations in the parent **before** child's launch are visible to the child.
- All global memory operations of the child are visible to the parent **after** the parent synchronizes on the child's completion.



# Why Multi-GPU?

- Having multiple CPU-GPU handshakes should suffice?



# Multiple Devices

- In general, a CPU may have different types of devices, with different compute capabilities.
- However, they all are nicely numbered from 0..N-1.
- *cudaSetDevice(i)*

What is wrong with this code from parallelization perspective?

```
cudaSetDevice(0);  
K1<<<...>>>();  
cudaMemcpy();  
cudaSetDevice(1);  
K2<<<...>>>();  
cudaMemcpy();
```

```
cudaSetDevice(0);  
K1<<<...>>>();  
cudaMemcpyAsync();  
cudaSetDevice(1);  
K2<<<...>>>();  
cudaMemcpyAsync();
```

# Multiple Devices

- `cudaGetDeviceCount(&c);`
  - Identify the number of devices.
- `cudaDeviceCanAccessPeer(&can, from, to);`
  - Can from device access to device?
- `cudaDeviceEnablePeerAccess(peer, ...);`
- While at the hardware level, the relation seems symmetric, the programming interface enforces asymmetry.
- Maximum 8 peer connections per device.
- Need 64 bit application.

# Enumerate Devices

```
int deviceCount;  
cudaGetDeviceCount(&deviceCount);  
int device;  
for (device = 0; device < deviceCount; ++device) {  
    cudaDeviceProp deviceProp;  
    cudaGetDeviceProperties(&deviceProp, device);  
    printf("Device %d has compute capability %d.%d.\n",  
          device, deviceProp.major, deviceProp.minor);  
}
```

# Warp Voting

- `__all(predicate);`
  - If all warp threads satisfy the predicate.
- `__any(predicate);`
  - If any warp threads satisfies the predicate.
- `__ballot(predicate);`
  - Which warp threads satisfy the predicate.
  - Generalizes `__all` and `__any`.

# Warp Voting

```
#include <stdio.h>
#include <cuda.h>

__global__ void K() {
    unsigned val = __all(threadIdx.x < 100);
    if (threadIdx.x % 32 == 0) printf("%X\n", val);
}

int main() {
    K<<<1, 128>>>();
    cudaDeviceSynchronize();

    return 0;
}
```

1  
1  
1  
0

# Warp Voting

```
#include <stdio.h>
#include <cuda.h>

__global__ void K() {
    unsigned val = __any(threadIdx.x < 100);
    if (threadIdx.x % 32 == 0) printf("%X\n", val);
}

int main() {
    K<<<1, 128>>>();
    cudaDeviceSynchronize();

    return 0;
}
```

1  
1  
1  
1

# Warp Voting

```
#include <stdio.h>
#include <cuda.h>

__global__ void K() {
    unsigned val = __ballot(threadIdx.x < 100);
    if (threadIdx.x % 32 == 0) printf("%X\n", val);
}

int main() {
    K<<<1, 128>>>();
    cudaDeviceSynchronize();

    return 0;
}
```

FFFFFFFF  
FFFFFFFF  
FFFFFFFF  
F



# Warp Voting

```
#include <stdio.h>
#include <cuda.h>

__global__ void K() {
    unsigned val = __ballot(threadIdx.x % 2 == 0);
    if (threadIdx.x % 32 == 0) printf("%X\n", val);
}

int main() {
    K<<<1, 128>>>();
    cudaDeviceSynchronize();

    return 0;
}
```

55555555  
55555555  
55555555  
55555555

# Warp Voting for atomics

- `if (condition) atomicInc(&counter, N);`
  - Executed by many threads in a warp, but not all.
  - The contention is high.
  - Can be optimized with `__ballot`.
- Leader collects warp-count.
  - `__ballot()` provides a mask; how do we count bits?
  - `__popc(mask)` returns the number of set bits.
  - `__ffs(mask)` returns the first set bit (from lsb).
- Leader performs a single `atomicAdd`.
  - Reduces contention.

# Warp Consolidation

Original code

```
if (condition) atomicInc(&counter, N);
```

Optimized code

```
unsigned mask = __ballot(condition);  
if (threadIdx.x % 32 == 0)  
    atomicAdd(&counter, __popc(mask));
```

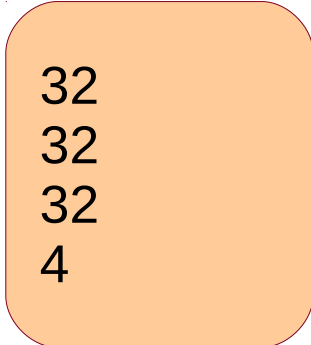
# Warp Voting for atomics

```
#include <stdio.h>
#include <cuda.h>

__global__ void K() {
    unsigned val = __ballot(threadIdx.x < 100);
    if (threadIdx.x % 32 == 0) printf("%d\n", __popc(val));
}

int main() {
    K<<<1, 128>>>();
    cudaDeviceSynchronize();

    return 0;
}
```



32  
32  
32  
4

# Conditional Warp Voting

- If a warp-voting function is executed within a conditional, some threads may be masked, and they would not participate in the voting.

```
if (threadIdx.x % 2 == 0) {  
    unsigned mask = __ballot(threadIdx.x < 100);  
    if (threadIdx.x % 32 == 0) printf("%d\n", __popc(mask));  
}
```

16  
16  
16  
2