

Multi-Dimensional arrays in CUDA

Matrix Squaring (CPU)

```
void squarecpu(unsigned *matrix, unsigned *result,  
               unsigned matrixsize /* = 32 */) {  
    for (unsigned ii = 0; ii < matrixsize; ++ii) {  
        for (unsigned jj = 0; jj < matrixsize; ++jj) {  
  
            for (unsigned kk = 0; kk < matrixsize; ++kk) {  
                result[ii * matrixsize + jj] +=  
                    matrix[ii * matrixsize + kk] * matrix[kk * matrixsize + jj];  
            }  
        }  
    }  
}
```

Matrix Squaring (GPU-Version 2)

```
square<<<N, N>>>(matrix, result, N);    // N = 32
```

```
__global__ void square(unsigned *matrix,  
                        unsigned *result,  
                        unsigned matrixsize) {  
  
    unsigned id = blockIdx.x * blockDim.x + threadIdx.x;  
    unsigned ii = id / matrixsize;  
    unsigned jj = id % matrixsize;  
    for (unsigned kk = 0; kk < matrixsize; ++kk) {  
        result[ii * matrixsize + jj] += matrix[ii * matrixsize + kk] *  
                                         matrix[kk * matrixsize + jj];  
    }  
}
```

Matrix Squaring (GPU-Version 3)

```
int main(){
    int i, j;
    int A[N][N];
    int B[N][N];
    dim3 BlockPerGrid(1, 1);
    dim3 ThreadPerBlock(N, N);

    // Initialize A to required values B to all zeros
    cudaMemcpyToSymbol(dA,A,size,0,cudaMemcpyDefault);
    cudaMemcpyToSymbol(dB,B,size,0,cudaMemcpyDefault);

    sqr <<< BlockPerGrid, ThreadPerBlock >>> (N);

    cudaMemcpyFromSymbol(A,dB,size,0,cudaMemcpyDefault);
    // Print the final results stored in A

    return 0;
}
```

Matrix Squaring (GPU-Version 3)

```
sqr<<<BlockPerGrid,ThreadPerBlock>>>(N); // N = 32
```

```
__device__ int dA[N][N];
__device__ int dB[N][N];

__global__ void sample(int matrixsize)
{
    int j = threadIdx.x + blockDim.x * blockIdx.x;
    int i = threadIdx.y + blockDim.y * blockIdx.y;

    //printf("%d %d\n",i, j);

    for(int k=0;k<matrixsize;k++)
        dB[i][j] += dA[i][k] * dA[k][j] ;
}
```

2D Minimum Algorithm

2	3	4	5
3	5	6	2
2	8	4	5
1	4	8	9



2	3	2
2	4	2
1	4	4

- 2D operations like this are found in many fundamental algorithms like Interpolation, Convolution, Filtering
- Applications in seismic processing, weather simulation, image processing, etc

2D Minimum Algorithm

- Each output element is the minimum of input elements within the window

$$O[i][j] = \min(A[i-1][j], A[i-1][j+1], A[i][j], A[i][j+1])$$

- The rows are in the range 0..N-1 and columns in 0..M-1
- The computation boundaries remains intact.
- Initialize all the values in A using rand() function.

2D Minimum Algorithm (CPU)

```
void init(int grid[][N], int prevgrid[][N])
{
    for(int i=0;i<N;i++){
        for(int j=0;j<N;j++){
            int val = rand() % 10;
            grid[i][j] = val;
            prevgrid[i][j] = val;
        }
    }
}

void compute(int grid[][N], int prevgrid[][N])
{
    // preform the computation
    for(int i=1;i<N;i++){
        for(int j=0;j<N-1;j++){
            grid[i][j] = min( min( prevgrid[i-1][j], prevgrid[i-1][j+1] ) ,
min( prevgrid[i][j], prevgrid[i][j+1] ) );
        }
    }
}
```


2D Minimum Algorithm (GPU)

```
void compute(int grid[][N], int prevgrid[][N])           // N = 32
{
    dim3 BlockPerGrid(1, 1);
    dim3 ThreadPerBlock(N, N);
    size_t size = N * N * sizeof(int);

    cudaMemcpyToSymbol(dgrid, grid, size, 0, cudaMemcpyDefault);
    cudaMemcpyToSymbol(dprevgrid, prevgrid, size, 0,
cudaMemcpyDefault);

    sample <<<BlockPerGrid, ThreadPerBlock>>> (N);

    cudaMemcpyFromSymbol(grid, dgrid, size, 0, cudaMemcpyDefault);
}
```

2D Minimum Algorithm (GPU)

```
sample <<<BlockPerGrid, ThreadPerBlock>>> (N);    // N = 32
```

```
__device__ int dgrid[N][N];
__device__ int dprevgrid[N][N];

__global__ void sample(int matrixsize){
    int j = threadIdx.x + blockDim.x * blockIdx.x;
    int i = threadIdx.y + blockDim.y * blockIdx.y;

    //printf("%d %d\n",i, j);

    if( i > 0 && j < N-1 )
        dgrid[i][j] = min( min( dprevgrid[i-1][j], dprevgrid[i-1]
[j+1] ) , min( dprevgrid[i][j], dprevgrid[i][j+1] ) );
}
```


Stencil Computation

- Design and implement a stencil computation wherein in each iteration, cells compute their values using the neighboring ones, according to the following formula:


$$A[i][j] += A[i-1][j] + A[i-1][j+1]$$

- The rows are in the range 0..N-1 and columns in 0..M-1
- The computation boundaries remains intact.
- Initialize all the values in A to all 1s.
- Your program should run for ten iterations.

Stencil Computation (in-place)



Input Matrix



Output Matrix

Stencil Computation (CPU)

```
// Perform the computation
for(int itr=0;itr<10;itr++){
```

```
    // perform the computation for the current iteration
    for(int i=1;i<N;i++){
        for(int j=0;j<N-1;j++){
            stencil[i][j] += prevstencil[i-1][j] + prevstencil[i-1][j+1];
        }
    }
```

```
    // copy the result to the current iteration to the prevstencil so that can
    be used for next iteration
```

```
    for(int i=0;i<N;i++){
        for(int j=0;j<N;j++){
            prevstencil[i][j] = stencil[i][j];
        }
    }

}
```

Stencil Computation (GPU)

```
dim3 BlockPerGrid(1, 1);
dim3 ThreadPerBlock(N, N);
size_t size = N * N * sizeof(int);

int itr;
// Perform the computation
for(itr=0;itr<10;itr++)
{
    cudaMemcpyToSymbol(dstencil, stencil, size, 0,
cudaMemcpyDefault);
    cudaMemcpyToSymbol(dprevstencil, prevstencil, size, 0,
cudaMemcpyDefault);
    sample <<<BlockPerGrid, ThreadPerBlock>>> (N);
    cudaMemcpyFromSymbol(prevstencil, dstencil, size, 0,
cudaMemcpyDefault);
    cudaMemcpyFromSymbol(stencil, dstencil, size, 0,
cudaMemcpyDefault);
}
```

Stencil Computation (GPU)

```
sample <<<BlockPerGrid, ThreadPerBlock>>> (N);    // N = 32
```

```
__device__ int dstencil[N][N];
__device__ int dprevstencil[N][N];

__global__ void sample(int matrixsize)
{
    int j = threadIdx.x + blockDim.x * blockIdx.x;
    int i = threadIdx.y + blockDim.y * blockIdx.y;

    //printf("%d %d\n",i, j);

    if( i > 0 && j < N-1 )
        dstencil[i][j] += dprevstencil[i-1][j] + dprevstencil[i-1][j+1];
}
```

Stencil Computation(Original)

- Design and implement a stencil computation wherein in each iteration, cells compute their values using the neighboring ones, according to the following formula:

$$A[i][j] += A[i-1][j] + A[i-1][j+1]$$

- The rows are in the range 0..N-1 and columns in 0..M-1
- The computation boundaries remains intact
- Initialize all the values in A to all 1s
- Your program should run for ten iterations

Stencil Computation (in-place)

- The updates should happen in-place inside the input matrix.


On CPU:

```
for(int itr=0;itr<10;itr++){  
    for(int i=N-1;i>=1;i--){  
        for(int j=0;j<N-1;j++){  
            stencil[i][j] += stencil[i-1][j] + stencil[i-1][j+1];  
        }  
    }  
}
```

Stencil Computation (in-place)

- The updates should happen in-place inside the input matrix


On GPU: `stencil[7][7]`



Stencil Computation (in-place)

- The updates should happen in-place inside the input matrix


On GPU: `stencil[7][7]`



Stencil Computation (in-place)

- The updates should happen in-place inside the input matrix


On GPU: `stencil[7][7]`



Stencil Computation (in-place)

- The updates should happen in-place inside the input matrix


On GPU: `stencil[7][7]`



Stencil Computation (in-place)

- The updates should happen in-place inside the input matrix

On GPU: `stencil[7][7]`



Stencil Computation (in-place)

```
size_t size = N * N * sizeof(int);
int itr;

// Perform the computation
for(itr=0;itr<10;itr++)
{
    cudaMemcpyToSymbol(dstencil, stencil, size, 0, cudaMemcpyDefault);

    for(int i=N-1;i>=1;i--)
    {
        cudaMemcpyToSymbol(dstencil[i], stencil[i], N, 0,
cudaMemcpyDefault);
        sample <<<1, N>>> (N,i);
        cudaMemcpyFromSymbol(stencil[i], dstencil[i], N, 0,
cudaMemcpyDefault);
    }

    cudaMemcpyFromSymbol(stencil, dstencil, size, 0,
cudaMemcpyDefault);
}
```

Stencil Computation (in-place)

```
sample <<<1, N>>> (N,i);           // N = 32, i = [1,N-1]
```

```
__device__ int dstencil[N][N];

__global__ void sample(int matrixsize, int i)
{
    int j = threadIdx.x + blockDim.x * blockIdx.x;

    //printf("%d %d\n",i, j);

    if( i > 0 && j < N-1 )
        dstencil[i][j] += dstencil[i-1][j] + dstencil[i-1][j+1];
}
```